


# Introduction


- Every new desktop/laptop is now equipped with a graphic processing unit (GPU).
- GPU = **Massively Parallel Architecture**.
- For most of their life, such GPUs are **idle**.
- General Purpose GPU applications:



Bioinformatics



Deep Learning



Numerical Analysis  
MathWorks MATLAB


# (Distributed) Constraint Optimization

- A *(D)COP* is a tuple  $\langle X, D, F, (A, \alpha) \rangle$ , where:
  - $X$  is a set of variables.
  - $D$  is a set of finite domains.
  - $F$  is a set of utility functions:  $f_i : \times_{x_j \in \text{scope}(f_i)} D_j \mapsto \mathbb{N} \cup \{0, -\infty\}$
  - $A$  is a set of agents, controlling the variables in  $X$ .
  - $\alpha$  maps variables to agents.
- **GOAL**: Find a utility maximal assignment.

$$\begin{aligned} \mathbf{x}^* &= \arg \max_{\mathbf{x}} \mathbf{F}(\mathbf{x}) \\ &= \arg \max_{\mathbf{x}} \sum_{f \in \mathbf{F}} f(\mathbf{x}|_{\text{scope}(f)}) \end{aligned}$$

# Graphical Processing Units (GPUs)

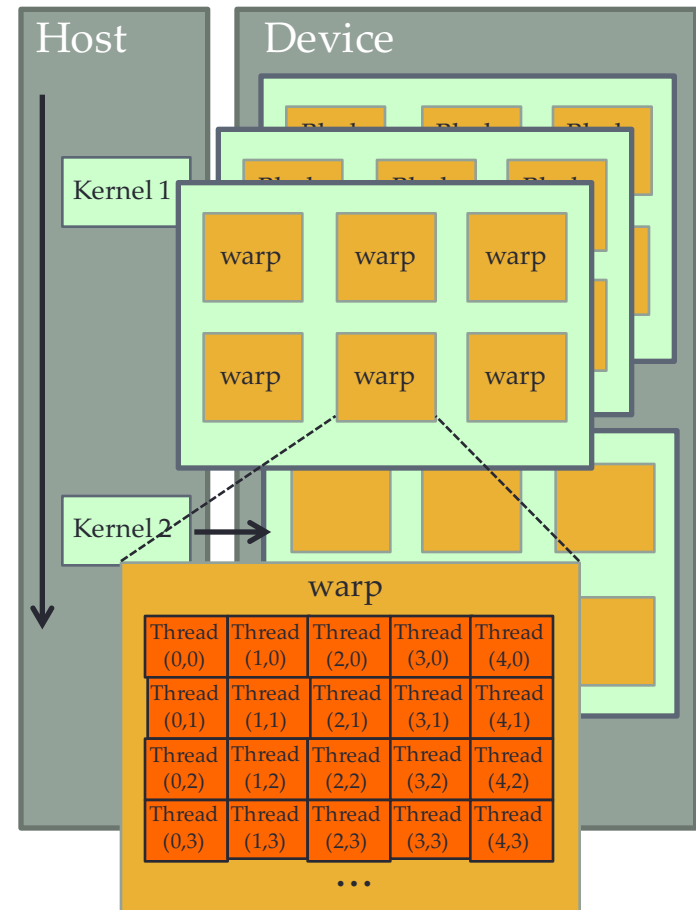
- A GPU is a massive parallel architecture:
  - **Thousands** of multi-threaded **computing cores**.
  - Very **high** **memory bandwidths**.
  - ~80% of transistors devoted to data processing rather than caching.



- **However:**
  - GPU cores are **slower** than CPU cores.
  - GPU memories have **different sizes** and **access times**.
  - GPU programming is more challenging and time consuming.

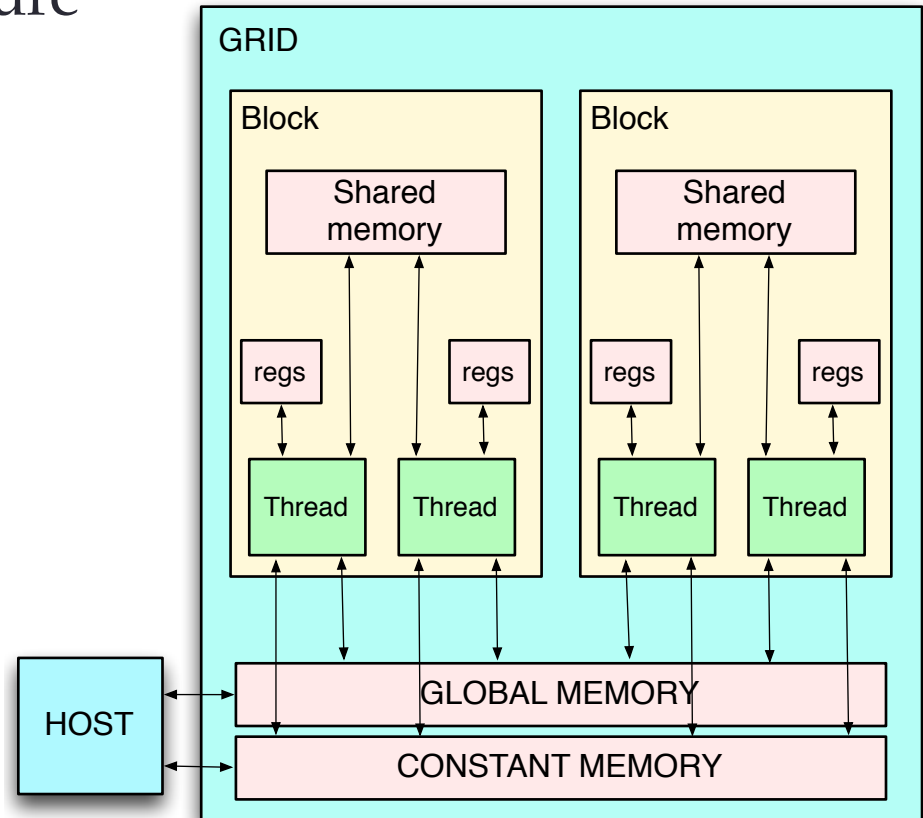
# Execution Model

- A **Thread** is the basic parallel unit.
- **Threads** are organized into a **Block**.
- Several **warps** are scheduled for the execution of a GPU function.
- Several **Streaming Multiprocessors**, (SM) scheduled in parallel.
- Single Instruction Multiple Thread (SIMT) parallel model.



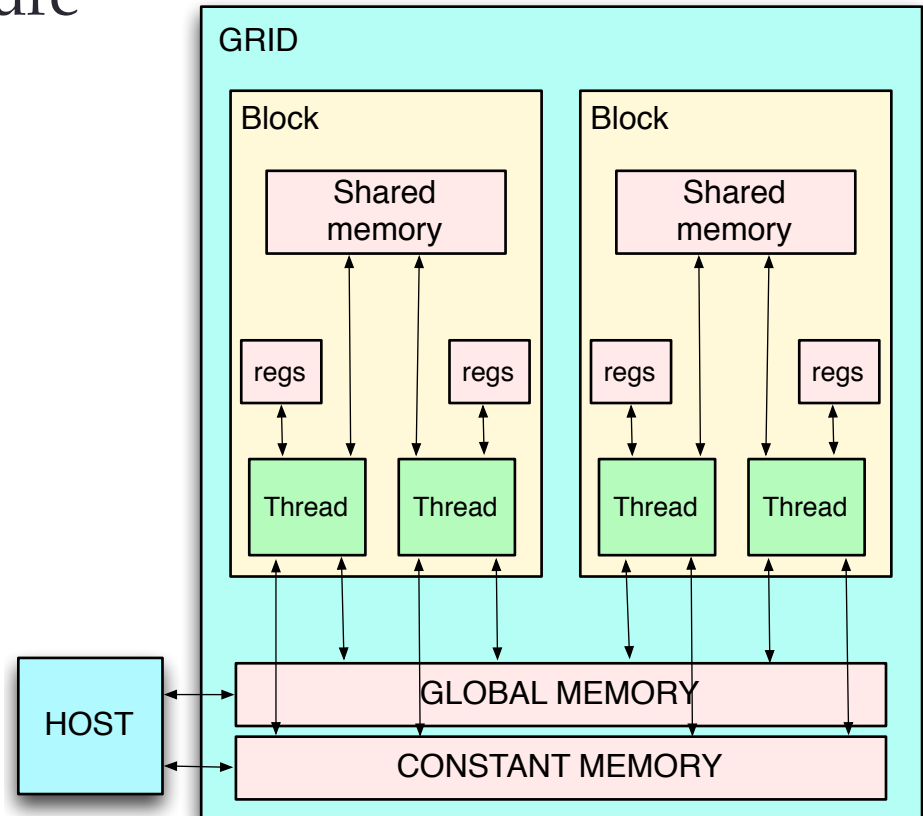
# Memory Hierarchy

- The GPU memory architecture is rather involved.
- Registers
- Shared memory
- Global memory




# Memory Hierarchy

- The GPU memory architecture is rather involved.
- **Registers**
  - Fastest;
  - Only accessible by a **thread**;
  - Lifetime of a thread.
- **Shared memory**
- **Global memory**



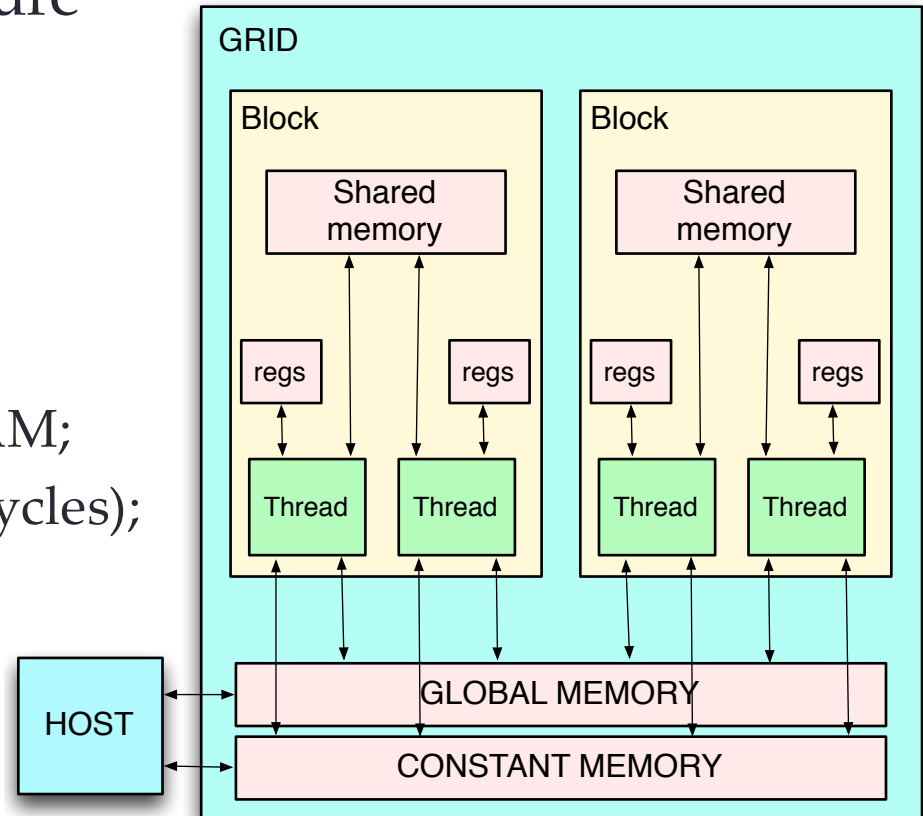
# Memory Hierarchy

- The GPU memory architecture is rather involved.
- **Registers**
- **Shared memory**
  - Extremely fast;
  - Highly parallel;
  - Restricted to a **block**.
- **Global memory**



# Memory Hierarchy

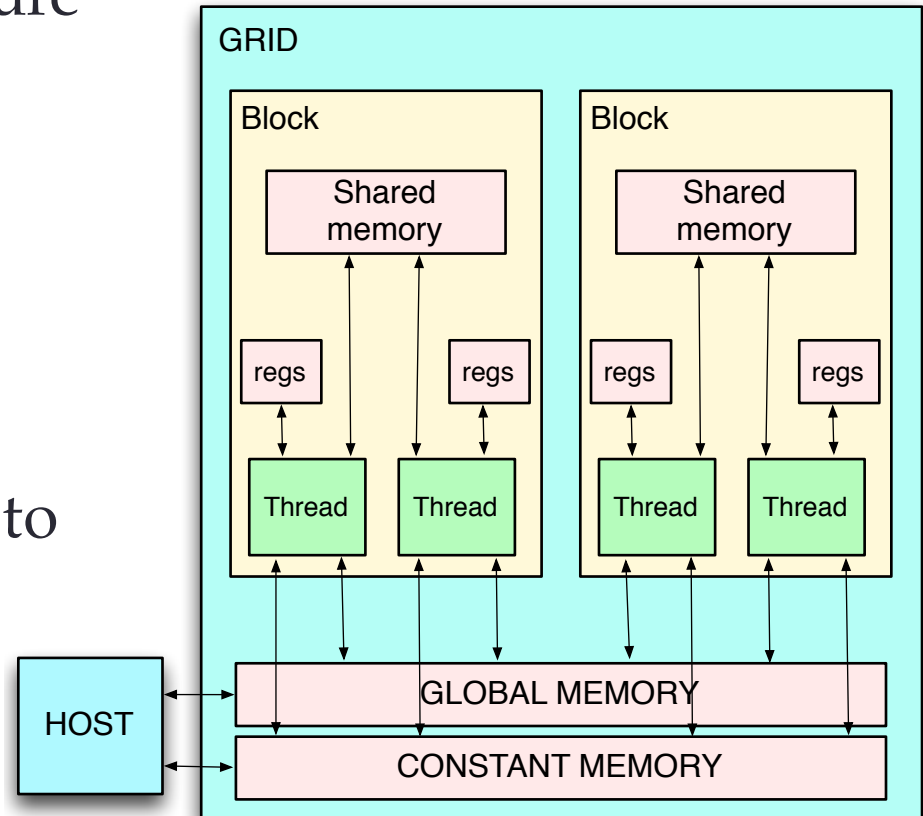
- The GPU memory architecture is rather involved.
- **Registers**
- **Shared memory**
- **Global memory**
  - Typically implemented in DRAM;
  - **High access latency** (400-800 cycles);
  - Potential of **traffic congestion**.





# Memory Hierarchy

- The GPU memory architecture is rather involved.
- **Registers**
- **Shared memory**
- **Global memory**
- **Challenge:** using memory effectively -- likely requires to redesign the algorithm.



# CUDA: Compute Unified Device Architecture

## Host



## Device



# CUDA: Compute Unified Device Architecture

Host



Device



```
cudaMalloc(&deviceV, sizeV);
```

```
cudaMemcpy(deviceV, hostV,  
            sizeV, ...)
```

data




Global Memory

# CUDA: Compute Unified Device Architecture

Host



Device



`cudaKernel<nThreads, nBlocks>()`

Kernel invocation



`cudaKernel()`



Global Memory

# CUDA: Compute Unified Device Architecture

Host



Device



`cudaMemcpy(hostV, deviceV,  
sizeV, ...)`

data



Global Memory

# Bucket Elimination and DPOP

- Dynamic Programming procedures to solve (D)COPs.
- Both procedures rely on the use of two operators:
- **Projection Operator:**  $\pi_{-x_i}(f_{ij})$

$x_i$	$x_j$	U	$\longrightarrow$	$x_j$	U
0	0	5		0	20
0	1	8		1	8
1	0	20			
1	1	3			

$f_{ij}$

- **Aggregation Operator:**  $f_{ij} + f_{ik}$

# Bucket Elimination and DPOP

- Dynamic Programming procedures to solve (D)COPs.
- Both procedures rely on the use of two operators:
- **Projection Operator:**  $\pi_{-xi}(f_{ij})$

$x_i$	$x_j$	U	$\Rightarrow \max(5, 20)$	$x_j$	U
0	0	5		0	20
0	1	8		1	8
1	0	20			
1	1	3			
$f_{ij}$					

- **Aggregation Operator:**  $f_{ij} + f_{ik}$

# Bucket Elimination and DPOP

- Dynamic Programming procedures to solve (D)COPs.
- Both procedures rely on the use of two operators:
- **Projection Operator:**  $\pi_{-x_i}(f_{ij})$

$x_i$	$x_j$	U		$x_j$	U
0	0	5		0	20
0	1	8		1	8
1	0	20	$\searrow$		
1	1	3	$\nearrow$		

$f_{ij}$

$\max(8, 3)$

- **Aggregation Operator:**  $f_{ij} + f_{ik}$



# Bucket Elimination and DPOP

- Dynamic Programming procedures to solve (D)COPs.
- Both procedures rely on the use of two operators:
- **Projection Operator:**  $\pi_{-xi}(f_{ij})$
- **Aggregation Operator:**  $f_{ij} + f_{ik}$

$x_i$	$x_j$	U
0	0	5
0	1	8
1	0	20
1	1	3



$x_i$	$x_k$	U
0	0	2
0	1	6
1	0	11
1	1	4


$f_{ij}$



$x_i$	$x_j$	$x_k$	U
0	0	0	7
0	0	1	11
0	1	0	10
0	1	1	14
...			
...			


# Bucket Elimination and DPOP

- Dynamic Programming procedures to solve (D)COPs.
- Both procedures rely on the use of two operators:
- **Projection Operator:**  $\pi_{-xi}(f_{ij})$
- **Aggregation Operator:**  $f_{ij} + f_{ik}$




# Bucket Elimination and DPOP

- Dynamic Programming procedures to solve (D)COPs.
- Both procedures rely on the use of two operators:
- **Projection Operator:**  $\pi_{-xi}(f_{ij})$
- **Aggregation Operator:**  $f_{ij} + f_{ik}$



# Bucket Elimination and DPOP

1. Imposes an ordering on the problem's variables.




$$X = \{x_1, x_2, x_3\}$$

$$F = \{f_{12}, f_{13}, f_{23}\}$$

# Bucket Elimination and DPOP

2. Selects the variable  $x_i$  with highest priority, and it creates a **bucket**:




$$B_i = \left\{ f_j \in \mathbf{F} \mid x_i \in \text{scope}(f_j) \wedge \right. \\ \left. i = \max\{k \mid x_k \in \text{scope}(f_j)\} \right\}$$

$$X = \{x_1, x_2, x_3\} \quad B_3 = \{f_{13}, f_{23}\}$$

$$F = \{f_{12}, f_{13}, f_{23}\}$$

# Bucket Elimination and DPOP

3. It computes a new utility function  $f_i'$  by **aggregating** the functions in  $B_i$  and **projecting** out  $x_i$



$$f_{13}$$

$x_1$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_{23}$$

$x_2$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_3'$$

$x_1$	$x_2$	U
0	0	$\max(5 + 5, 8 + 8) = 16$

(x<sub>3</sub>)  
1


$$f_3' = \pi_{-x_3} (f_{13} + f_{23})$$

$$X = \{x_1, x_2, x_3\} \quad B_3 = \{f_{13}, f_{23}\}$$

$$F = \{f_{12}, f_{13}, f_{23}\}$$

# Bucket Elimination and DPOP

3. It computes a new utility function  $f_i'$  by **aggregating** the functions in  $B_i$  and **projecting** out  $x_i$



$$f_{13}$$

$x_1$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_{23}$$

$x_2$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_3'$$

$x_1$	$x_2$	U	$(x_3)$
0	0	16	1
0	1	$\max(5 + 20, 8 + 3) = 25$	0


$$f_3' = \pi_{-x_3} (f_{13} + f_{23})$$

$$X = \{x_1, x_2, x_3\} \quad B_3 = \{f_{13}, f_{23}\}$$

$$F = \{f_{12}, f_{13}, f_{23}\}$$

# Bucket Elimination and DPOP

3. It computes a new utility function  $f_i'$  by **aggregating** the functions in  $B_i$  and **projecting** out  $x_i$



$$f_{13}$$

$x_1$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_{23}$$

$x_2$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_3'$$

$x_1$	$x_2$	U	$(x_3)$
0	0	16	1
0	1	25	0
1	0	<b>max(20 + 5, 3 + 8) = 25</b>	0

$$f_3' = \pi_{-x_3} (f_{13} + f_{23})$$


$$X = \{x_1, x_2, x_3\} \quad B_3 = \{f_{13}, f_{23}\}$$

$$F = \{f_{12}, f_{13}, f_{23}\}$$



# Bucket Elimination and DPOP

3. It computes a new utility function  $f_i'$  by **aggregating** the functions in  $B_i$  and **projecting** out  $x_i$



$$f_{13}$$

$x_1$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_{23}$$

$x_2$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$$f_3'$$

$x_1$	$x_2$	U	$(x_3)$
0	0	16	1
0	1	25	0
1	0	25	0
1	0	$\max(20 + 20, 3 + 3) = 40$	0


$$f_3' = \pi_{-x_3} (f_{13} + f_{23})$$

$$X = \{x_1, x_2, x_3\} \quad B_3 = \{f_{13}, f_{23}\}$$

$$F = \{f_{12}, f_{13}, f_{23}\}$$

# Bucket Elimination and DPOP

4. It updates the set of variables:  $\mathbf{X} \leftarrow \mathbf{X} \setminus \{x_i\}$
5. It updates the set of functions:  $\mathbf{F} \leftarrow (\mathbf{F} \cup \{f'_i\}) \setminus B_i$




$$\mathbf{X} = \{x_1, x_2\}$$

$$B_2 = \{f_{13}, f_3'\}$$

$$\mathbf{F} = \{f_{12}, f_3'\}$$

# Bucket Elimination and DPOP


Repeat...



$$X = \{x_1, x_2\}$$


$$F = \{f_{12}, f_3'\}$$

- **DPOP** is a distributed version of BE.
- It operates on a **Pseudotree** ordering of the constraint graph.



# GPU-(D)BE

- BE and DPOP complexity:  $O(d^{w^*})$ .  
 $d$  = max. domain size;  $w^*$  = induced width of the constraint graph.
- Can the projection and aggregator operators be executed in parallel?
- Do they fit the SIMT parallel model?



$x_1$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3


$x_2$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

$x_1$	$x_2$	U
0	0	$\max(5 + 5, 8 + 8) = 16$
0	1	$\max(5 + 20, 8 + 3) = 25$
1	0	$\max(20 + 5, 3 + 8) = 25$
1	0	$\max(20 + 20, 3 + 3) = 40$

$$f_3' = \pi_{-x3} (f_{13} + f_{23})$$

# GPU-(D)BE

- BE and DPOP complexity:  $O(d^{w^*})$ .  
 $d$  = max. domain size;  $w^*$  = induced width of the constraint graph.
- Can the projection and aggregator operators be executed in parallel?
- Do they fit the SIMT parallel model?
- **Obs.:** The computation of each row of the Utility tables is independent from the computation of other rows.



$x_1$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3


$x_2$	$x_3$	U
0	0	5
0	1	8
1	0	20
1	1	3

	$x_1$	$x_2$	U
→	0	0	$\max(5 + 5, 8 + 8) = 16$
→	0	1	$\max(5 + 20, 8 + 3) = 25$
→	1	0	$\max(20 + 5, 3 + 8) = 25$
→	1	1	$\max(20 + 20, 3 + 3) = 40$

$$f_3' = \pi_{-x_3}(f_{13} + f_{23})$$

# Algorithm design and data structure

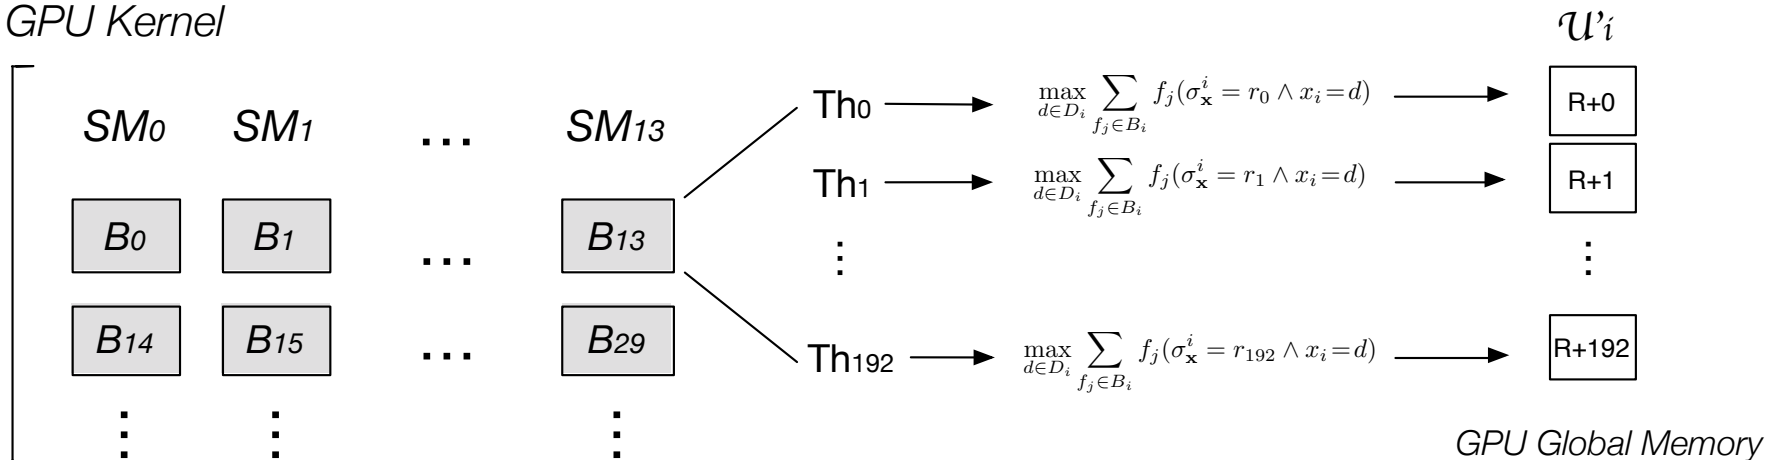
- Limit the amount of host-device **data transfers**.
  - **Static Entities**: require a **single** data transaction.
    - Variables; Domains; Utility functions; Constraint Graph.
  - **Dynamic Entities**: might require **multiple** data transactions.
    - Utility tables.
- Minimize the accesses to the **global memory**.
  - Padding Utility Tables' rows; Perfect hashing.
- Ensure data accesses are **coalesced**.
  - Mono-dimensional array organization;



# Parallel Projection and Aggregation

- Mapping between the  $f_i'$  table rows and the *CUDA* blocks:
  - Each **thread** in a block is associated to the computations of **one permutation** of values in  $scope(f_i')$ .
  - 1 **block** =  $64k$  threads ( $1 \leq k \leq 16$ ).
    - $k$  depends on the architecture and it is chosen so to maximize the number of threads that can be scheduled concurrently.
- **Obs.:** Max number of parallel  $f_i'$  table rows is  $M = |SM| 64k$ 
  - In our experiments,  $|SMs| = 14$  and  $k = 3$ . Thus  **$M = 2688$** .

GPU Kernel




# Experimental Results

- $|D_i| = 5$ ;
- $p_2 = 0.5$ ;
- timeout = 300s

## Regular Grid Networks

- CPU: 2.3GHz, 128 GB RAM
- GPU: 14 SMs, 837MHz.



## Speedup

avg. max.: 125.1x

avg. min.: 42.6x

- Similar trends at increasing  $|D_i|$ .
- Similar trends for DPOP vs GPU-DBE

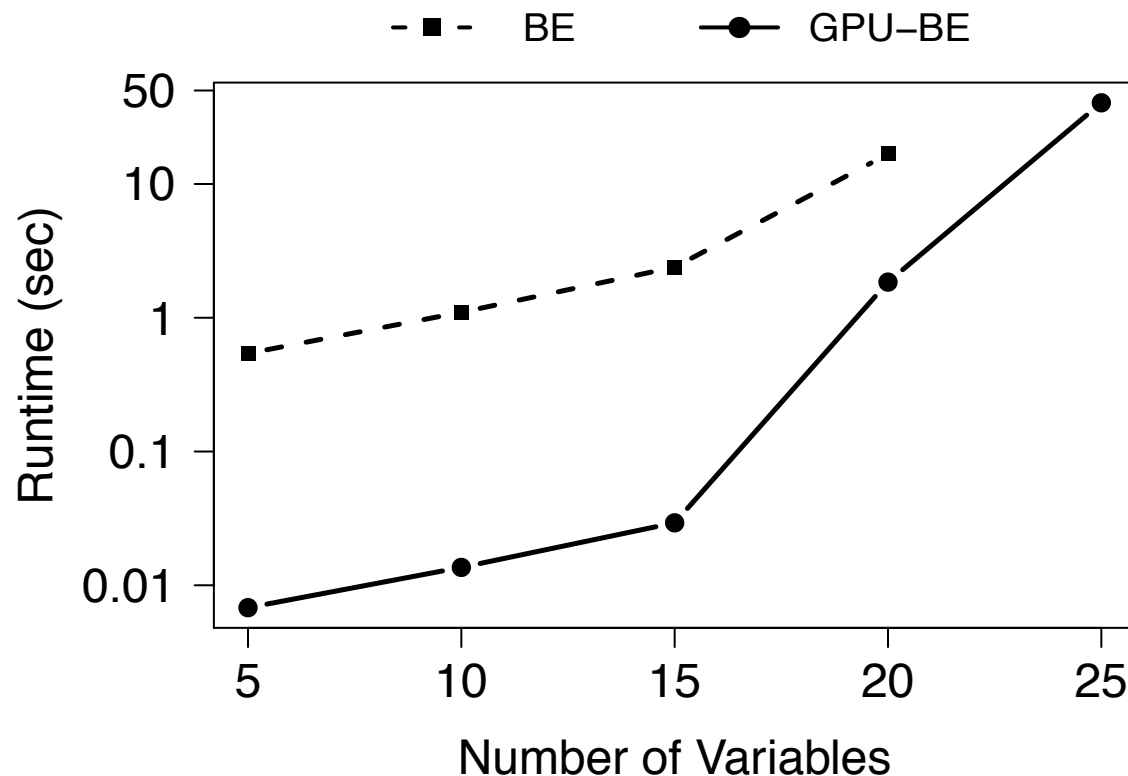


# Experimental Results

- $|D_i| = 5$ ;
- $p_2 = 0.5$ ;  $p_1 = 0.3$ ;
- timeout = 300s

## Random Networks

- CPU: 2.3GHz, 128 GB RAM
- GPU: 14 SMs, 837MHz.



## Speedup

avg. max.: 69.3x

avg. min.: 16.1x


- Similar trends at increasing  $|D_i|$ .
- Similar trends for DPOP vs GPU-DBE

# Experimental Results

- $|D_i| = 5$ ;
- $p_2 = 0.5$ ;
- timeout = 300s

## Scale Free Networks

- CPU: 2.3GHz, 128 GB RAM
- GPU: 14 SMs, 837MHz.



## Speedup

avg. max.: 34.9x

avg. min.: 9.5x

- Similar trends at increasing  $|D_i|$ .
- Similar trends for DPOP vs GPU-DBE

# Lesson Learned #1


- The  $f_i'$  table size increases exponentially with  $w^*$ .
- Limited GPU global memory (2GB).
- $f_i'$  table +  $B_i$  tables, to be used in the aggregation operation, might exceed global memory capacity!
- Partition  $f_i'$  computations in **multiple chunks**.
- Alternates GPU and CPU to compute  $f_i'$ .
  - **GPU**: Aggregates the functions in  $B_i$  excluding those which do not fit in the global memory.
  - **CPU**: Aggregates the other functions in  $B_i$ ;  
Projects out the variable  $x_i$ .

# Experimental Results

- $|A| = 10$ ;  $|D_i| = 5$ ;
- $p_2 = 0.5$ ;
- $\text{timeout} = 300\text{s}$

## Random Networks

- CPU: 2.3GHz, 128 GB RAM
- GPU: 14 SMs, 837MHz.



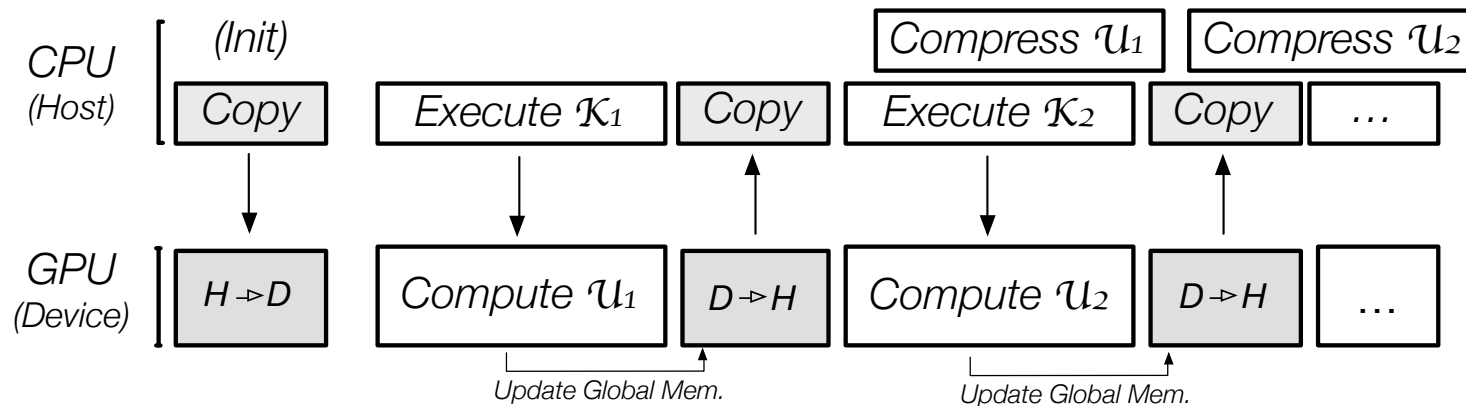
## Phase Transition

$$p_1 = 0.4$$

- small  $p_1$  correspond to smaller  $w^*$

# Lesson Learned #2

- Host and Device **concurrency**.
- Possible when the  $f_i'$  tables are computed in **chunks**.
- It may hide host-device data transfers as byproduct.



# Discussion

- Exploiting the **integration of CPU and GPU** is a key factor to obtain competitive solver performance.
- How to determine good tradeoffs of such integration?
- **GPU:**
  - Repeated, non memory intensive operations;
  - Operations requiring regular memory access;
- **CPU:**
  - Memory intensive operations;


# Conclusions

- Exploit GPU-style parallelism from DP-based (D)COPs resolution methods.
- GPU-(D)BE: Exploits GPUs to parallelizes the **aggregation** and **projection** operators.
- Observed different speedup, ranging from 34.9 to 125.1, based on several network topologies.
- Discussed several possible optimization techniques.
- **FUTURE WORK:**
  - Exploit GPUs in DP-based propagators.
  - Investigate GPUs in higher form of consistency.

# Exploiting GPUs in Solving (Distributed) Constraint Optimization Problems with Dynamic Programming

F. Fioretto, T. Le, E. Pontelli, W. Yeoh, T. Son

## Thank You!



**Ferdinando Fioretto**

New Mexico State University,  
University of Udine

**Email:** [ffiorett@cs.nmsu.edu](mailto:ffiorett@cs.nmsu.edu)

**Web:** [www.cs.nmsu.edu/~ffiorett](http://www.cs.nmsu.edu/~ffiorett)